

Credit underwriting ML project

Fred Serfati - 07/23/2024



Plan

1. Introduction: Credit underwriting
2. EDA
3. Data Preprocessing & Feature Engineering
4. Data Visualization
5. Modeling
6. Results and Feature importances
7. Conclusions
8. Q&A



Introduction: Credit underwriting

- Determine the creditworthiness of an applicant, by identifying whether they should be given a loan in the future
- Helps financial institutions manage risks and ensure profitability
- Accurate credit underwriting decisions:
 - Protect lenders from financial losses
 - Support borrowers in accessing fair credit opportunities



Modelisation using Machine Learning

- **Supervised problem:** access to outcome/label of loan
- **Binary classification:** categorical target with 2 categories: good (non-default) & bad (default)



Exploratory Data Analysis



Data preparation

- 2 datasets:
 - **Application** (644 records) for every customer that has been given a loan in a 6 month period
 - **Loan** (1266 records) for the outcome of those loans: good (i.e., non default) or bad (default)
- **Merged** Loan data (target) to Application data (variables) using a **left join**
- Final dataset: only 631 records: **tiny dataset!**




Dataset presentation

- 631 records, 32 columns
- **Target:** flgGood, i.e., whether a loan defaulted ('Bad', 0) or not ('Good', 1)
- Different types of variables:
 - **Numerical:** e.g., amount_requested, monthly_rent_amount
 - **Categorical:**
 - **Ordinal:** e.g., email_duration, residence_duration
 - **Nominal:** e.g., bank_account_direct_deposit, how_use_money
- (Stratified) train test split:
 - **80% (488) for train/validation**, used for model training, fine-tuning and model selection
 - **20% (122) for test**, used for predictions (**nothing else!**)



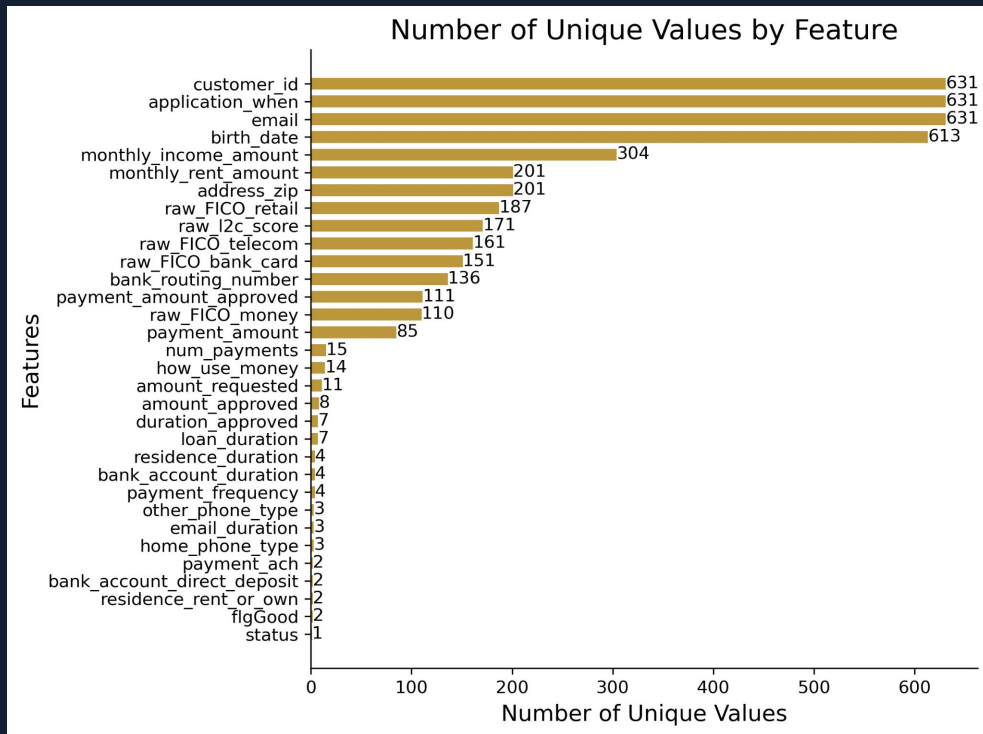
Dealing with missing values



- Other_phone_type: dropped column
- 3 other features: dropped rows with null values (very few)
- Tried imputation with mean for payment_amount_approved but let to poorer performances



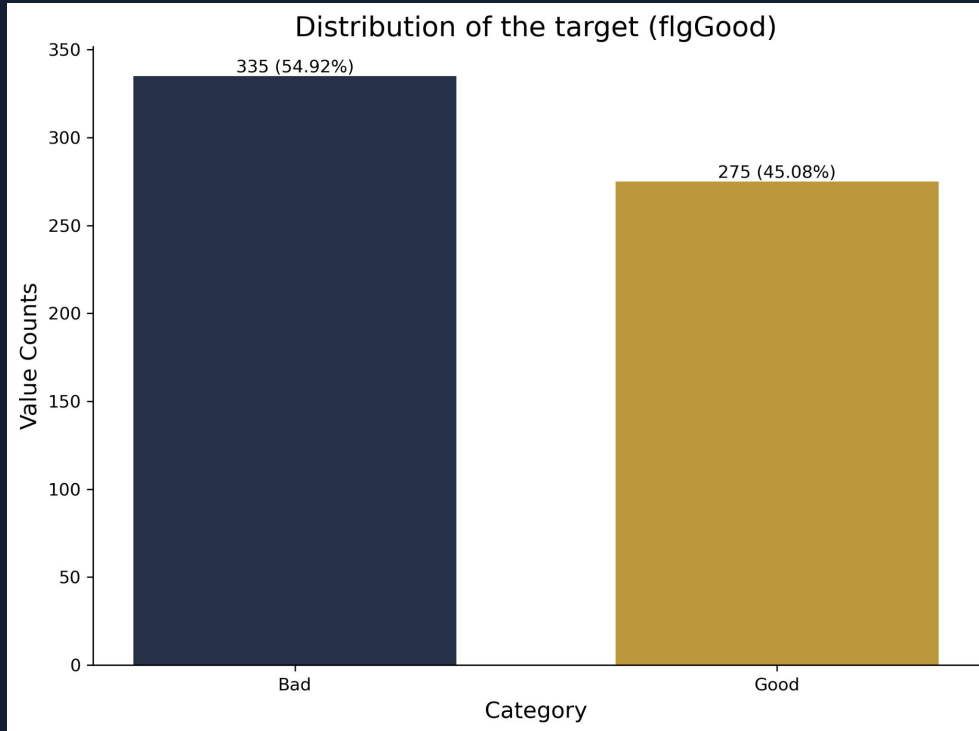
Unique values



- Customer_id, application_when and email: dropped or transformed (1 unique value per application)
- Status: dropped (1 unique value, because all loans in this dataset were approved; i.e., 0 variance)



A slightly imbalanced dataset



- **45% of good loans vs 55% of bad loans: slight imbalance**
- Still very limited, so no need for advanced sampling methods (such as SMOTE)
- However, **choice of metrics** still **very important** (can't use accuracy)



Data Preprocessing & Feature Engineering



Data preprocessing

- **Numerical features:**
 - Standard scaling
 - No need for missing values imputation: already removed all missing values
- **Nominal variables:** one-hot-encoding
- **Ordinal variables:** ordinal encoding



Feature engineering

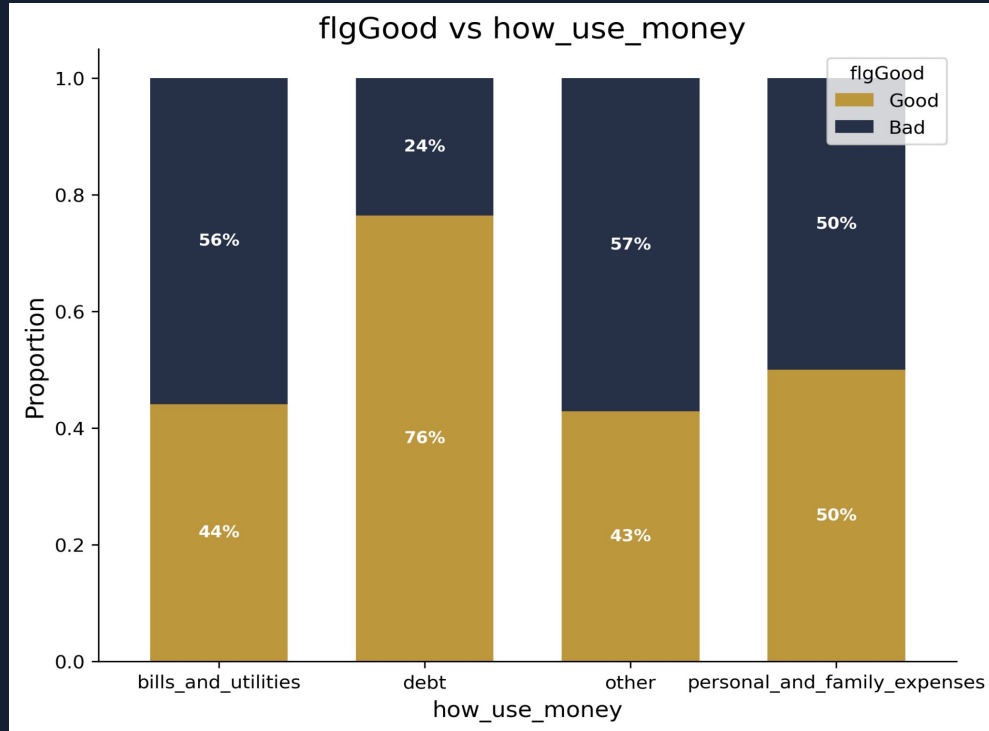
- Age
- Year, month and day of week of application
- How_use_money: categories regroupment
- Zip code: only kept 3 first digits to reduce dimensionality
- Average FICO score (money, retail, telecom, bank card)
- Potentially useful ratios, e.g.:
 - Debt-to-income (DTI) and approved DTI
 - Approved-to-requested-amount: proportion of the requested amount that was granted
 - Approved-to-requested-loan-duration: proportion of the requested duration that was granted



Data Visualization



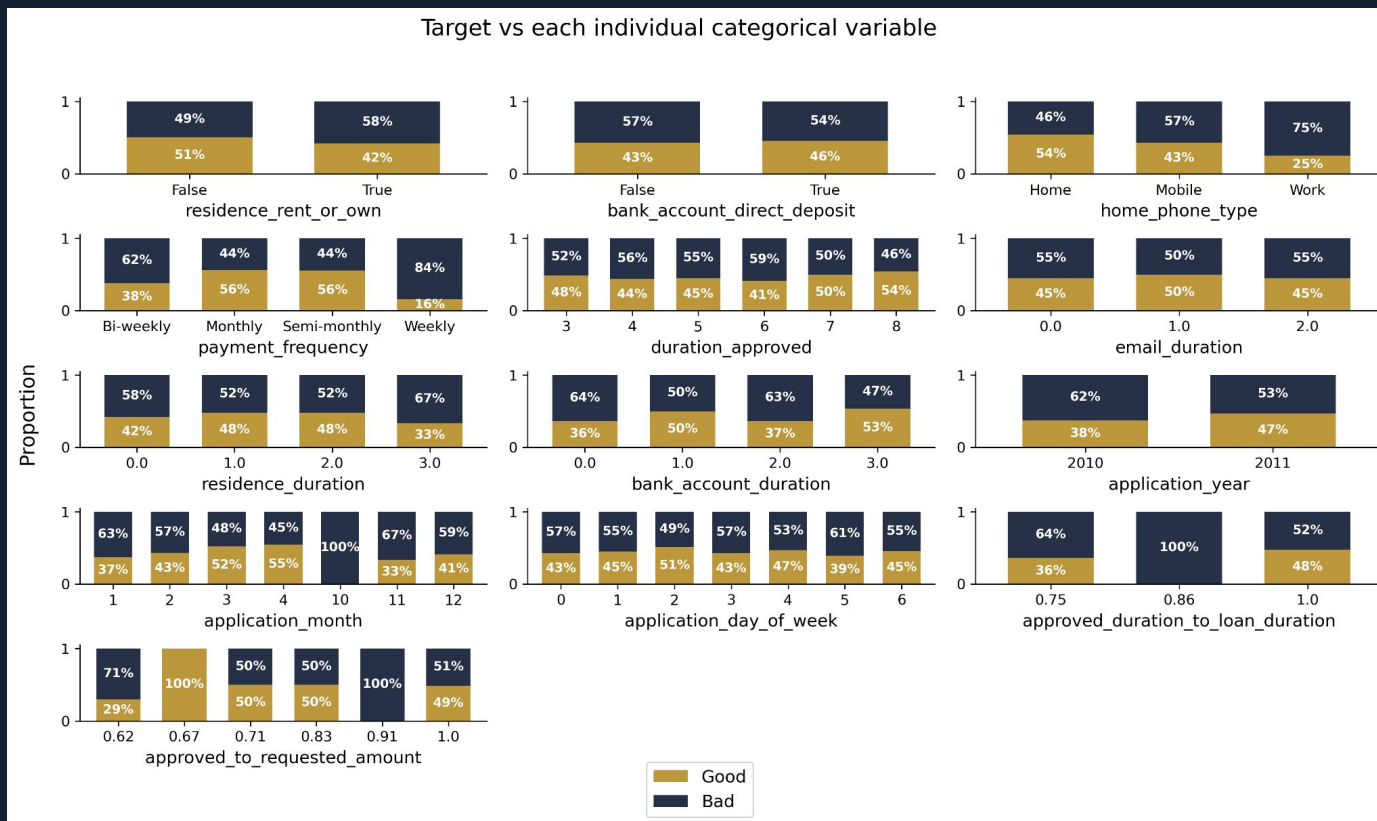
Relationship between target and categorical features (1/2)



- 76% of loans used for debt reimbursement were good
- Only 50% for personal and family expenses

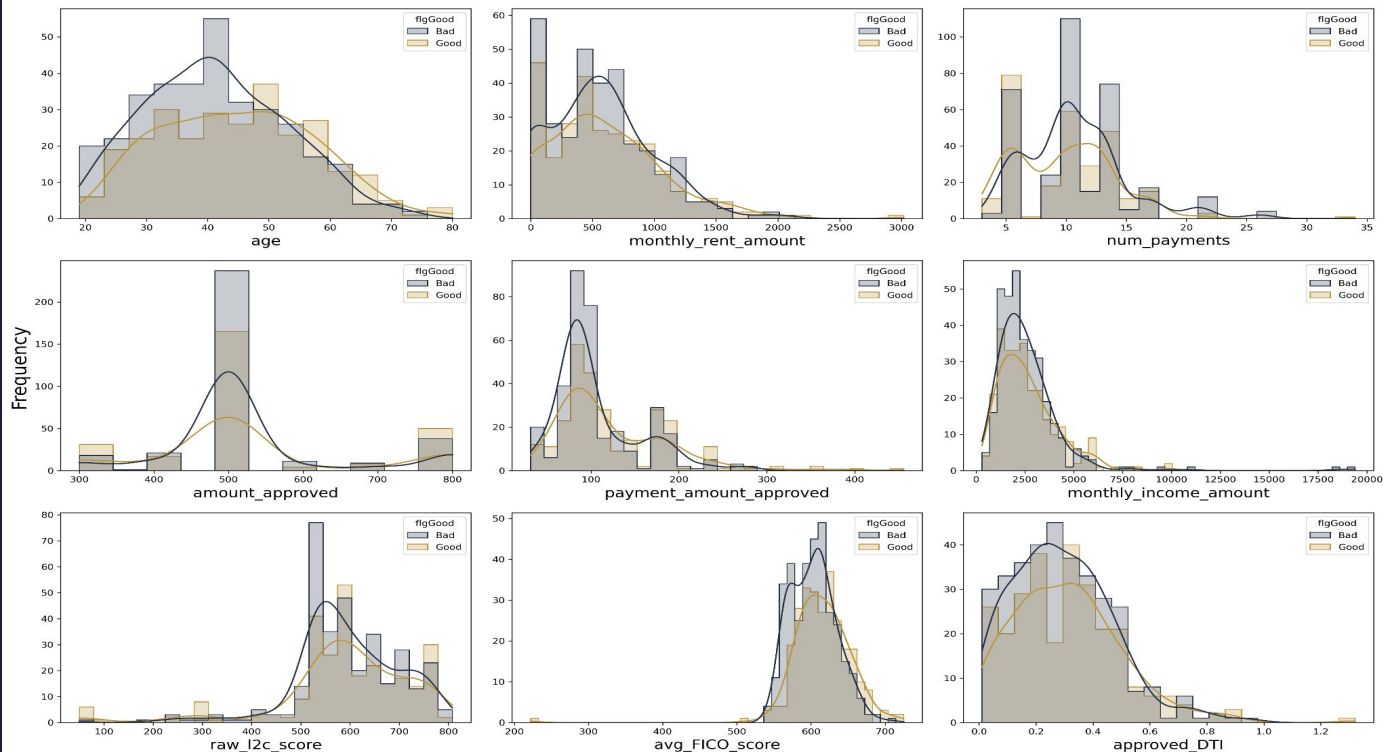


Relationship between target and categorical features (2/2)




Relationship between target and numerical features

Histogram of Numerical Variables by flgGood




Pairwise relationships between numerical features



- Strong correlation between:
 - Requested & approved duration
 - Requested & approved amount
 - Requested & approved payment amount
 - Loan duration & number of payments
 - All FICO scores
 - ...
- Moderate correlation between:
 - Num_payments & approved_duration
 - Num_payment & amount_approved
 - ...




Pairwise relationships between ordinal features



Nothing remarkable here: no “correlation” between any pair of ordinal features



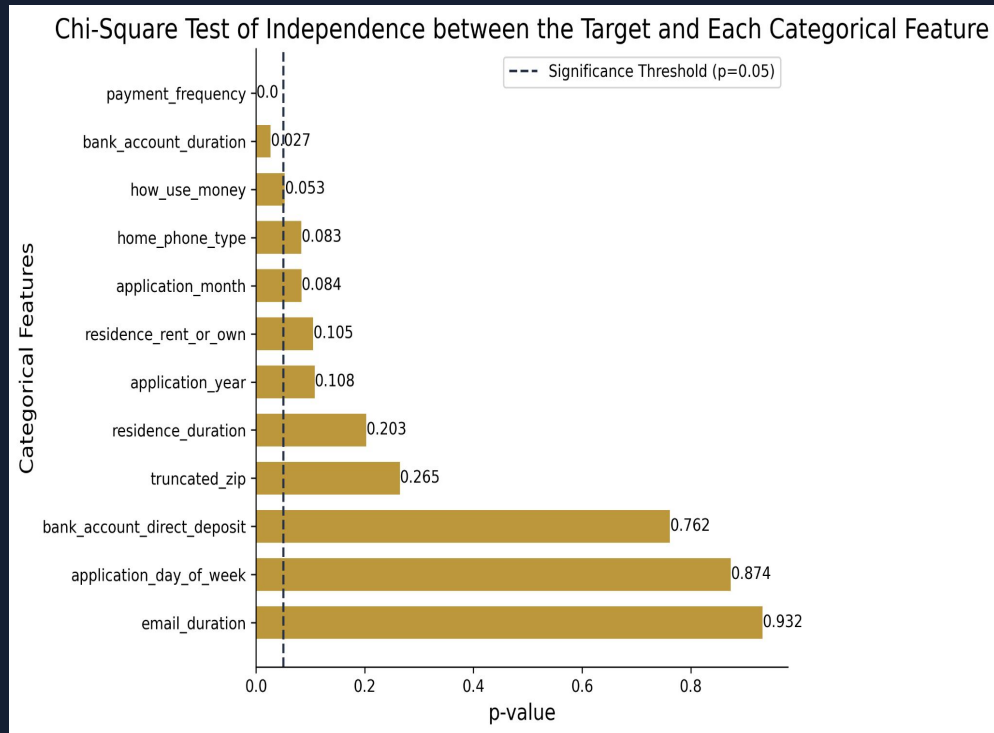
Pairwise relationships between nominal features



- Strong relationship between application year & month:
 - Probably because data from 2010 only cover October to December while data from 2011 only cover January April
- Moderate relationship between application month (and year) & how_use_money



Chi-square tests of independence



- H_0 : ind. vs H_1 : non-ind.
- Payment_frequency, bank_account_duration and how_use_money:
 - **Reject H_0** : target not independent of those variables
- All other categorical variables:
 - **Fail to reject H_0** : not enough evidence to conclude variable not independent from target
- Careful with interpretation because of multiple testing problem



Modeling



Why we can't use accuracy

- (Slight) imbalance problem, so biased towards majority class ('Bad')
- More importantly, in credit underwriting, we don't only care about overall correctness, but also about:
 - **False positives** (granting a credit that will default: **money loss**)
 - **False negatives** (missing a good loan: **missed opportunity**)



Metrics used

- **Balanced accuracy:** weighted version of accuracy, taking into account class imbalance
- **Precision:** proportion of true positive (good loans) out of all loans predicted as good
 - The higher, the more we can trust that if a loan is predicted as good, it will really be good
- **Recall (TPR):** proportion of true positives (good loans) out of all actual good loans
 - The higher, the more good loans are identified by the model
- **F1-score:** balance between precision and recall
 - Useful here, as we care both about precision and recall
- **AUC-ROC:**
 - The higher, the better the model is able to distinguish between good and bad loans
- **AU-PRC:**
 - The higher, the better the model performs with the positive (good) class



Model evaluation: stratified 5-fold cross-validation

- Why **cross-validation** instead of a **simple train/valid split**?
 - **Very small training dataset** (less than 500 records):
 - High variance in metrics from one split to another: not trustworthy
 - To avoid bad surprises once model in production, need to robustly evaluate generalization error: **cross-validation**
 - The higher the number of folds, the more robust (but computationally more expensive)
- Why **stratified**?
 - Same proportion of Good vs Bad loans in each fold (compare apples to apples)



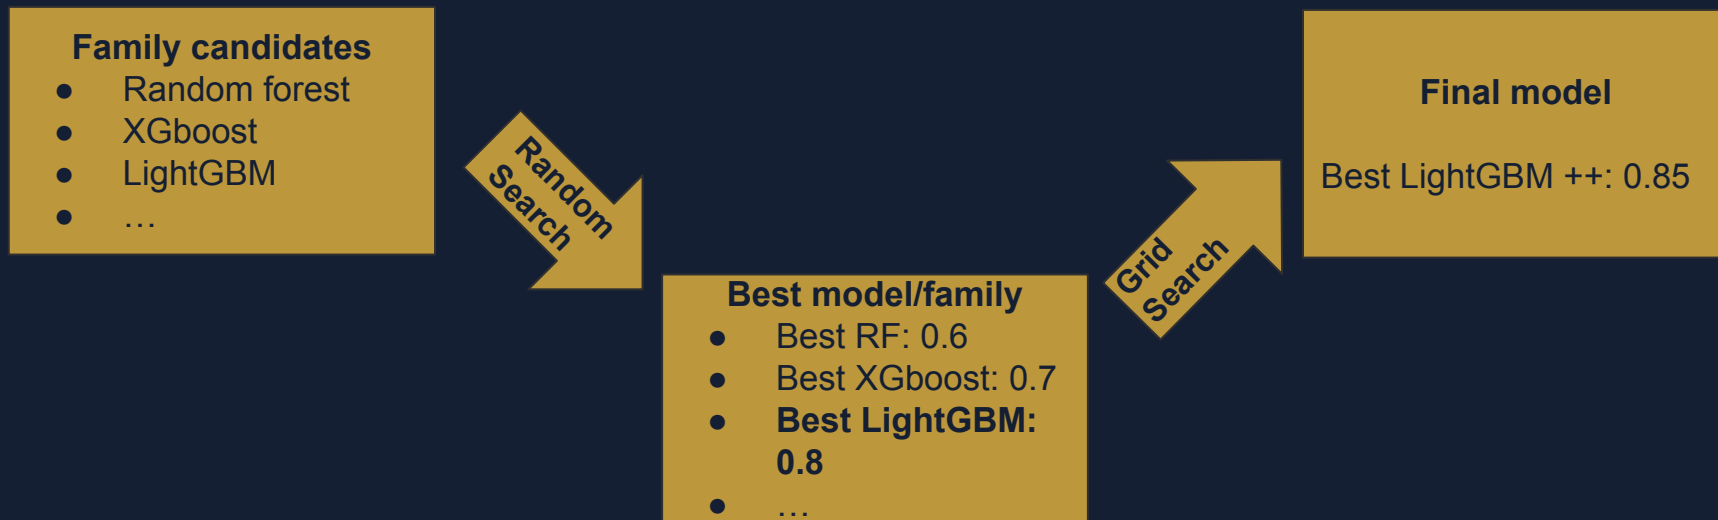
Model selection: random search then grid search

- For each model:
 - Narrow down the best hyperparameter region using a 100-iteration randomized search
- Select the best model according to the metrics *
- Define a refined and with lower cardinality hyperparameter grid around the parameters of the best model
- Run a grid search on the model with the previous grid to fine-tune it even further
- Best of both worlds:
 - efficiency of randomized search
 - accuracy of grid search

* What if different metrics lead to different best models? Fortunately not the case here, as there was always a consensus between metrics.



Model selection: random search then grid search



Machine learning models explored


- Dummy classifier (always predicts the majority class): serves as a **baseline**
- Logistic regression
- Tree based models:
 - Decision Tree
 - Random Forest
 - Gradient boosting:
 - XGBoost
 - LightGBM
 - CatBoost



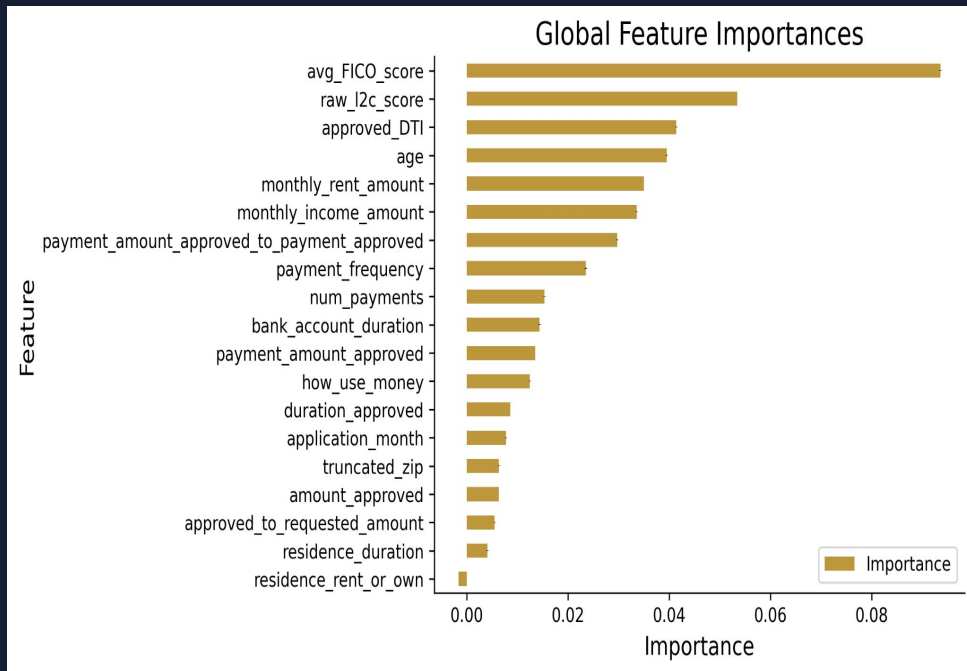
Results & Feature Importances



Results



Global Feature importances



- Method: **Permutation importance**
- **Relative** importances (not absolute)
- **Global** importances: most important features **overall**, not locally
- Most important features:
 - FICO score, L2C score, DTI, age, monthly rent amount, monthly income amount
- Least important features:
 - Residence rent or own, residence duration, approved to requested amount
- Careful with interpretation because of **correlated features** and **poor performances**



Final predictions (extract)

	predicted_probability	predicted_label	true_label
50	0.611983	1	1
51	0.595950	0	0
52	0.821817	0	0
53	0.694141	1	1
54	0.702674	0	0
55	0.659986	0	1
56	0.731026	1	1
57	0.672661	1	1
58	0.653270	1	1
59	0.720309	1	0



Conclusions



Challenges & Limitations

- Biggest challenge: **making the most out of such a small dataset**
- Lack of feature description:
 - Some variables were not self-explanatory, and thus I had to assume several things (e.g., residence_rent_or_own, or difference between payment amount and payment amount approved)



Recommendations & Next steps

- **Collect more data** and re-train the models
 - If no additional data are available, use data augmentation techniques to generate synthetic data
- **Discuss with domain experts** to gain knowledge about feature importances in credit underwriting
- Try other feature engineering techniques (e.g., cyclical encoding of time-related features)
- Use other feature importances methods, such as:
 - L2 logistic regression coefficients
 - Local methods for explaining individual predictions: **SHAP, LIME**



Thank you for your attention!
Questions?

